

ESE 345 Final Project Report

Alan George (114484206) and Matthew Leo (114238245)

Table of Contents:

Goals	2
Multimedia Unit Block Diagram	3
Design Procedure	3
Register File:	3
Instruction Buffer:	4
IF/ID Register:	5
ID/EX Register:	6
Data Forwarding Unit:	7
Forwarding Multiplexor:	8
Multimedia ALU:	9
EX/WB Register (Writeback):	9
Pipe Processor:	10
Testbench:	11
4-Stage Pipeline	13
Data Forwarding	15
Instruction Execution	17
LI	17
Signed Integer Multiply-Add Low with Saturation	17
Signed Integer Multiply-Add High with Saturation	18
Signed Integer Multiply-Subtract Low with Saturation	18
Signed Integer Multiply-Subtract High with Saturation	18
Signed Long Integer Multiply-Add Low with Saturation	19
Signed Long Integer Multiply-Add High with Saturation	19
Signed Long Integer Multiply-Subtract Low with Saturation	20
Signed Long Integer Multiply-Subtract High with Saturation	20
NOP	20
SHRHI	21
AU	21
CNT1H	21
AHS	22
OR	22
BCW	23
MAXWS	23
MINWS	23
MLHU	24
MLHSS	25
AND	25
ROTW	26

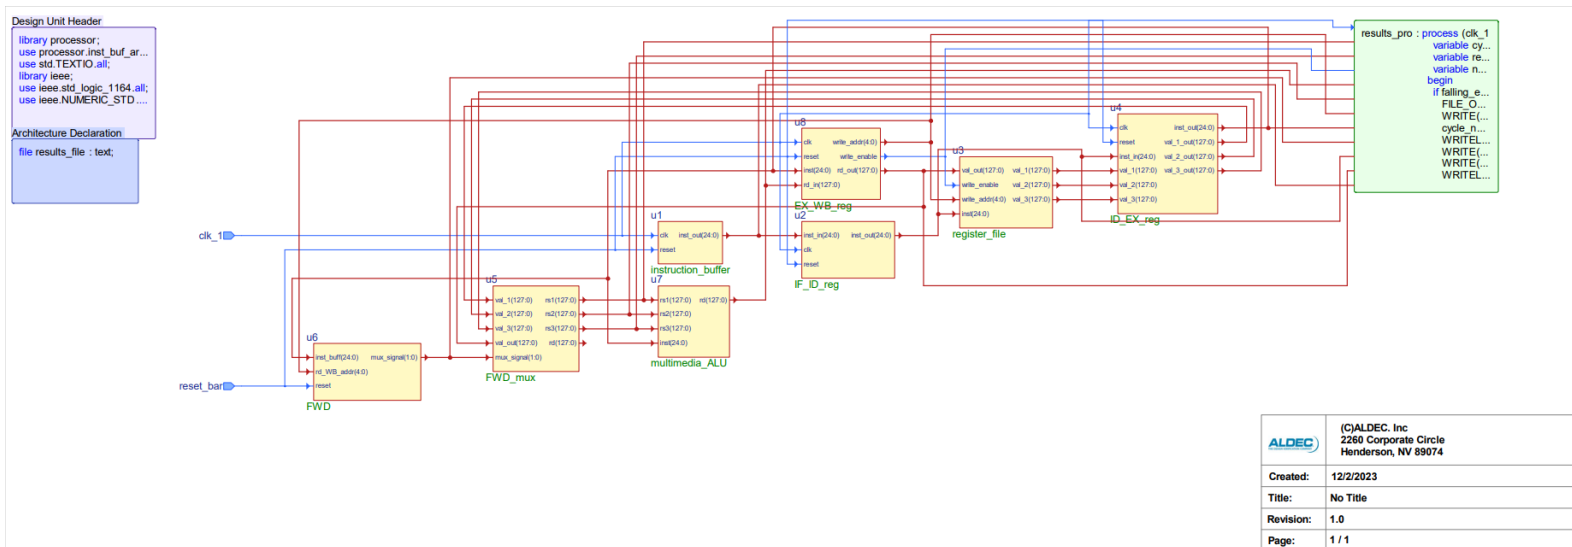
SFWU	27
SFHS	27
Simulation Results	27
Conclusion	31

Goals

The goal of this project was to create a 4-stage pipelined multimedia unit that included Instruction Fetch, Instruction Decode, Execution, and Write Back stages. The system needs to be capable of data forwarding to prevent data hazards and limit the need for pipeline stalls. The Register File in the Instruction Decode stage needs to be capable of storing 32 128-bit registers that can be written to and stored. The Instruction Buffer is expected to store up to 64 25-bit instructions that will be inputted into the pipeline each cycle. The multimedia ALU is able to process standard instructions such as AND and OR with multiple registers, along with also computing SIMD instructions that are similar to those found in the Sony Cell SPU and Intel SSE architectures.

Each component in the processor should also be written as a behavioral model in VHDL along with one structurally-written VHDL file that connects all of the components together with relevant signals. Our project also needs to take in a text file that has the binary machine code for all 64 instructions, and place them into the Instruction Buffer before reading in any instructions. We also need to produce a “Results File” that represents the status of the pipeline at each stage, by printing the instruction being processed in each stage of the pipeline, any relevant input operands, outputs, and control signals from each stage. Finally, our system should be efficient and able to execute each instruction with the minimal amount of cycles required.

Multimedia Unit Block Diagram



Design Procedure

Register File:

After the first part of the project, we decided to start with the register file. The register file was stated to be capable of holding 32 128-bit registers so we designed an array to hold all 32 registers.

architecture behavioral of register_file is

```
type reg_file is array (0 to 31) of std_logic_vector(127 downto 0);
signal sig_reg_file : reg_file := (others => (others => '0'));
```

begin

These registers would be accessed using the register addresses held in the 25-bit instruction. To differentiate between the R3, R4, and Load immediate, we made sure to use the codes given at the beginning of the instruction this way we could avoid mixing up immediates, opcodes, and rs3.

```

if inst(24) = '0' then
    val_1 <= sig_reg_file(to_integer(unsigned(inst(4 downto 0))));
    -- li instruction
    -- rd
    -- *** val_1 => rs1, val_1 => rd (in port_map to ALU)
else
    -- reading from register file
    val_1 <= sig_reg_file(to_integer(unsigned(inst(9 downto 5))));
    val_2 <= sig_reg_file(to_integer(unsigned(inst(14 downto 10))));
    val_3 <= sig_reg_file(to_integer(unsigned(inst(19 downto 15))));
    -- rs1
    -- rs2
    -- rs3
end if;

if inst(9 downto 5) = write_addr then
    val_1 <= val_out;
elsif inst(14 downto 10) = write_addr then
    val_2 <= val_out;
elsif inst(19 downto 15) = write_addr and inst(24 downto 23) = "10" then -- R4 inst only
    val_3 <= val_out;
end if;

```

Since we can have up to 3 register addresses stored in an instruction at a time, we created 3 outputs: *val_1*, *val_2*, and *val_3*. Using the addresses stored in the instruction we can access the stored data in the registers in the array and assign their 128-bit values to our outputs such as in our code shown above. (outputs shown below)

```

entity register_file is
    port(
        val_out : in STD_LOGIC_VECTOR(127 downto 0);
        write_enable : in std_logic;
        write_addr : in std_logic_vector(4 downto 0);
        -- value that we are writing to a register
        -- writing binary signal
        -- address of write-back register (rd)

        inst : in std_logic_vector(24 downto 0);
        -- IF/ID Reg (will pass inst to ALU)

        val_1 : out STD_LOGIC_VECTOR(127 downto 0);
        val_2 : out STD_LOGIC_VECTOR(127 downto 0);
        val_3 : out STD_LOGIC_VECTOR(127 downto 0);
        -- val from rs1
        -- val from rs2
        -- val from rs3
    );
end register_file;

```

Lastly, the register file is supposed to incorporate write back features so we can adjust the values stored in the array. For that reason, we designed two more inputs: *write_enable* and *write_addr*. As the name suggests, *write_enable* signifies when the system would write back to the register file and *write_addr* stores the 5-bit address of rd in the instruction so we know where to store the updated values from the writeback stage.

Instruction Buffer:

The next part we designed was the instruction buffer. The instruction buffer is loaded with 64 25-bit instructions that are given by the testbench and continuously fed into the system each clock cycle. To implement this we used a package created by the testbench to transfer instructions over to the instruction buffer. The instructions are stored in a 64 wide array each storing a 25-bit *std_logic_vector*.

```

package inst_buf_array is
    type inst_buf is array (0 to 63) of std_logic_vector(24 downto 0);
    shared_variable sig_inst_buf : inst_buf;
    -- inst_buff array input
    -- All 64 instructions in instruction buffer
end package;

```

The buffer has a *clk* input which makes sure that the instructions only get loaded each cycle.

```

entity instruction_buffer is
    port(
        clk : in std_logic;
        reset : in std_logic;

        inst_out : out std_logic_vector(24 downto 0)
    );
end instruction_buffer;

begin

    instruction_buffer : process(clk) is
        variable PC : integer := 0;
        -- variable cycle_num : integer := 1;
        -- variable results_line : line;
    begin

        if reset = '0' then
            PC := 0;
        else
            if rising_edge(clk) then

                if PC > 63 then
                    std.env.finish;
                end if;

                inst_out <= sig_inst_buf(PC);
                PC := PC + 1;

                -- if PC > 63 then
                --     null;
                -- else
                --     inst_out <= sig_inst_buf(PC);
                --     PC := PC + 1;
                -- end if;

            end if;
        end if;
    end process;
end;

```

On each rising edge we load one of the 64 instructions into our output called inst_out and send it to the IF/ID reg.

IF/ID Register:

The IF/ID register is basically a point that stores the instruction value till the clk increments and allows all stages to proceed, therefore, it is designed with 1 input, 1 output, and a clk input.

```

library IEEE;
use IEEE.std_logic_1164.all;

entity IF_ID_reg is
    port(
        inst_in : in STD_LOGIC_VECTOR(24 downto 0);
        clk : in STD_LOGIC;
        reset : in std_logic;

        inst_out : out STD_LOGIC_VECTOR(24 downto 0)
    );
end IF_ID_reg;

```

Each *rising_edge(clk)* the system will allow an instruction to move into the next stage.

```

architecture behavioral of IF_ID_reg is
begin
    if_id_buffer : process(clk) is
    begin
        if reset = '1' then
            if (rising_edge(clk)) then
                inst_out <= inst_in;
            end if;
        end if;
    end process if_id_buffer;
end behavioral;

```

ID/EX Register:

The ID/EX register holds the values being output from the Register File and holds them for a cycle till the next *rising_edge(clk)*. Since the Register File has 4 outputs, we have 4 inputs to store and then output

```

entity ID_EX_reg is
    port(
        clk : in std_logic;

        reset : in std_logic;

        inst_in : in STD_LOGIC_VECTOR(24 downto 0);
        val_1 : in STD_LOGIC_VECTOR(127 downto 0);
        val_2 : in STD_LOGIC_VECTOR(127 downto 0);
        val_3 : in STD_LOGIC_VECTOR(127 downto 0);

        inst_out : out STD_LOGIC_VECTOR(24 downto 0);
        val_1_out : out STD_LOGIC_VECTOR(127 downto 0);
        val_2_out : out STD_LOGIC_VECTOR(127 downto 0);
        val_3_out : out STD_LOGIC_VECTOR(127 downto 0)
    );
end ID_EX_reg;

```

```

architecture behavioral of ID_EX_reg is
--  file results_file : text;

begin
    id_ex_buffer : process(clk) is
        variable results_line : line;

    begin
        if reset = '1' then
            if (rising_edge(clk)) then
                inst_out <= inst_in;
                val_1_out <= val_1;
                val_2_out <= val_2;
                val_3_out <= val_3;

--                file_open(results_file, "../results.txt", append_mode);

--                write(results_line, "ID: ");
--                write(results_line, inst_in);
--                writeline(results_file, results_line);

--                write(results_line, inst_in);
--                writeline(results_file, results_line);

            end if;
        end if;
    end process id_ex_buffer;

end behavioral;

```

The instruction is also stored and passed every stage as a means of keeping track of each register address easily at the cost of efficiency.

Data Forwarding Unit:

At the Writeback stage, to ensure registers are given their most recent values we send out a writeback address to the data forwarding unit called *rd_WB_addr*.

```

entity FWD is
    port(
        inst_buff : in STD_LOGIC_VECTOR(24 downto 0);      -- coming from ID/EXE buffer
        rd_WB_addr : in STD_LOGIC_VECTOR(4 downto 0);      -- coming from EXE/WB buffer

        reset : in std_logic;

        mux_signal : out STD_LOGIC_VECTOR(1 downto 0)      -- towards Forwarding Muxes
    );
end FWD;

```

The forwarding unit compares the write back address with the incoming instruction's addresses to the ALU, these are stored in *inst_buff*. The forwarding unit checks first what type of instruction we are sending to the ALU (R3, R4, load immediate) and then checks rd, rs1, rs2, and rs3

depending on the type of instruction and then sends to the Forwarding Mux a *mux_signal* that tells it whether rd, rs1, rs2, or rs3 is being forwarded.

```

FWD : process(inst_buff, rd_WB_addr, reset) is
begin
    if reset = '0' or inst_buff = inst_UUU then
        mux_signal <= "00";
        rd_WB_addr <= "00000";
    else
        if inst_buff(24) = '0' then
            if inst_buff(4 downto 0) = rd_WB_addr then -- li instruction
                mux_signal <= "01";
            else
                mux_signal <= "00";
            end if;
        elsif inst_buff(9 downto 5) = rd_WB_addr then
            mux_signal <= "01";
        elsif inst_buff(14 downto 10) = rd_WB_addr then
            mux_signal <= "10";
        elsif inst_buff(19 downto 15) = rd_WB_addr and inst_buff(24 downto 23) = "10" then --R4 instruction data forwarding
            mux_signal <= "11";
        else
            mux_signal <= "00";
        end if;
    end if;
end process FWD;

```

Forwarding Multiplexor:

The Forwarding Mux takes in the 3 outputs from the ID/EX register, 1 output from the Writeback stage called *val_out*, and the *mux_signal* from the forwarding unit to determine whether or not registers need to be replaced by their most recent values in the Writeback stage. If the *mux_signal* is a "01" rs1 needs to be updated, if it's "10" then rs2 needs to be updated, if the signal is "11" then rs3 gets updated, else the registers are not being written back to. In the event of a Load immediate instruction, rs1 will also hold the value for the rd specified in the load immediate instruction.

```

entity FWD_mux is
port(
    val_1 : in STD_LOGIC_VECTOR(127 downto 0);
    val_2 : in STD_LOGIC_VECTOR(127 downto 0);
    val_3 : in STD_LOGIC_VECTOR(127 downto 0);
    val_out : in STD_LOGIC_VECTOR(127 downto 0) := (others => '0');

    mux_signal : in STD_LOGIC_VECTOR(1 downto 0) := "00";

    rs1 : out STD_LOGIC_VECTOR(127 downto 0);
    rs2 : out STD_LOGIC_VECTOR(127 downto 0);
    rs3 : out STD_LOGIC_VECTOR(127 downto 0)
);
end FWD_mux;

```



```

architecture behavioral of FWD_mux is
begin
    FWD_mux : process(mux_signal, val_out) is
    begin
        if mux_signal = "01" then
            rs1 <= val_out;
            rs2 <= val_2;
            rs3 <= val_3;
        elsif mux_signal = "10" then
            rs1 <= val_1;
            rs2 <= val_out;
            rs3 <= val_3;
        elsif mux_signal = "11" then
            rs1 <= val_1;
            rs2 <= val_2;
            rs3 <= val_out;
        else
            rs1 <= val_1;
            rs2 <= val_2;
            rs3 <= val_3;
        end if;
    end process FWD_mux;
end

```

Multimedia ALU:

As it was designed in the previous part of the project, the ALU takes in rs1, rs2, and rs3 from the forwarding mux and the instruction is still continuously passed from each stage in the pipeline.

```

entity multimedia_ALU is
port(
    rs1 : in std_logic_vector(127 downto 0) := (others => '0');
    rs2 : in std_logic_vector(127 downto 0);
    rs3 : in std_logic_vector(127 downto 0);
    inst: in std_logic_vector(24 downto 0);

    rd : out std_logic_vector(127 downto 0)
);
end multimedia_ALU;

```

After doing internal computations for each specific R3, R4 instructions and load immediate we send the result to rd to which is outputted to the EX/WB register.

EX/WB Register (Writeback):

In the EX/WB register we pass in the rd from the execute stage and the 25-bit instruction that we've been passing each stage.

```

entity EX_WB_reg is
    port(
        clk : in std_logic;

        reset : in std_logic;

        inst : in STD_LOGIC_VECTOR(24 downto 0);
        rd_in : in STD_LOGIC_VECTOR(127 downto 0);

        write_addr : out STD_LOGIC_VECTOR(4 downto 0);
        write_enable : out STD_LOGIC;
        rd_out : out STD_LOGIC_VECTOR(127 downto 0) := (others => '0')
    );
end EX_WB_reg;

```

We use the 25-bit instruction to determine if we have a nop or not. This is because if we have a “nop” we must make sure the *write_enable* is set to 0 otherwise it is always 1 to write back to the Register File and update our registers. Using the last 5 bits of the instruction we can store rd’s address in the write_address so we know which register in the register file we must update.

```

    elsif (rising_edge(clk)) then
        if inst(24 downto 23) = "11" and inst(18 downto 15) = "0000" then
            write_enable <= '0';
            -- nop case

        elsif inst = inst_UUU then
            write_enable <= '0';

        else
            write_enable <= '1';
            rd_out <= rd_in;
            write_addr <= inst(4 downto 0);
        end if;
    end if;
    -- everything else

```

After a rising_edge(clk) signifying the next cycle, we can send all our required write back information to the forwarding unit and Register File. *rd_out* stores the actual value in the register and *write_addr* stores the location of that register in the array located in the Register File.

Pipe Processor:

This is our structural architecture for the pipeline that links the system together.

```

-- instruction buffer
u1 : entity instruction_buffer port map (clk => clk_1, reset => reset_bar, inst_out => inst_output_1); -- output comes from a signal

-- IF/ID register
u2 : entity IF_ID_reg port map (inst_in => inst_output_1, clk => clk_1, reset => reset_bar, inst_out => inst_output_2);

-- register file
u3 : entity register_file port map (val_out => val_out_2, write_enable => write_enable_WB, write_addr => write_addr, inst => inst_output_2, val_1 => rs1, val_2 => rs2, val_3 => rs3);

-- ID/EX register
u4 : entity ID_EX_reg port map (clk => clk_1, reset => reset_bar, inst_in => inst_output_2, val_1 => rs1, val_2 => rs2, val_3 => rs3, inst_out => inst_output_3, val_1_out => rs1_id_ex, val_2_out => rs2_id_ex);

-- forwarding muxes
u5 : entity FWD_mux port map (val_1 => rs1_id_ex, val_2 => rs2_id_ex, val_3 => rs3_id_ex, val_out => val_out_2, mux_signal => mux_signal, rs1 => rs1_FWD_mux, rs2 => rs2_FWD_mux, rs3 => rs3_FWD_mux);

-- data forwarding unit
u6 : entity FWD port map (inst_buff => inst_output_3, rd_WB_addr => write_addr, reset => reset_bar, mux_signal => mux_signal);

-- ALU
u7 : entity multimedia_ALU port map (rs1 => rs1_FWD_mux, rs2 => rs2_FWD_mux, rs3 => rs3_FWD_mux, inst => inst_output_3, rd => rd_ALU);

-- EX/WB register
u8 : entity EX_WB_reg port map (clk => clk_1, reset => reset_bar, inst => inst_output_3, rd_in => rd_ALU, write_addr => write_addr, write_enable => write_enable_WB, rd_out => val_out_2);

```

This also handles our results file by using the *write()* and *writeline()* functions to list the status of the pipeline for each cycle during the program execution using the “clk” like the stages.

```

if falling_edge(clk_1) then
    file_open(results_file, "../results.txt", append_mode);

    --report to results file
    write(results_line, "Cycle #: " & integer'image(cycle_num));
    cycle_num := cycle_num + 1;
    writeline(results_file, results_line);

    write(results_line, "IF: ");
    write(results_line, inst_output_1);
    writeline(results_file, results_line);
--    writeline(results_file, new_line);

    write(results_line, "ID: ");
    write(results_line, inst_output_2);
    writeline(results_file, results_line);
--    writeline(results_file, new_line);

    write(results_line, "EX: ");
    write(results_line, inst_output_3);
    writeline(results_file, results_line);

    write(results_line, "    rs1: ");
    write(results_line, rs1_FWD_mux);
    writeline(results_file, results_line);

    write(results_line, "    rs2: ");
    write(results_line, rs2_FWD_mux);
    writeline(results_file, results_line);

    write(results_line, "    rs3: ");
    write(results_line, rs3_FWD_mux);
    writeline(results_file, results_line);

    write(results_line, "    rd: ");
    write(results_line, rd_ALU);
    writeline(results_file, results_line);

    write(results_line, "    Mux Signal (00 -> no forwarding): ");

```

Testbench:

The testbench basically manages the pipeline's cycles by setting the clock period.

```

library ieee;
use ieee.std_logic_1164.all;
use ieee.numeric_std.all;
use work.inst_buf_array.all;
use work.all;

entity testbench is
end testbench;

architecture behavioral of testbench is

    signal clk : std_logic := '0';      -- clk signal
    constant period : time := 1 us;

    signal reset : std_logic;            -- reset signal

begin

    processor_UUT : entity processor
    port map (
        clk_1 => clk,
        reset_bar => reset
    );

    reset <= '0', '1' after 2 * period;    -- reset signal

    clock: process                        -- system clock
    begin

        read_inst;                        -- read a new instruction from the input

        for i in 0 to (2**16) loop
            wait for period/2;
            clk <= not clk;
        end loop;
        std.env.finish;
    end process;

end behavioral;

```

The testbench also holds the package filled with the instructions which need to be inputted into the instruction buffer, all 64 25-bit instructions. It gets these instructions from the input.txt which it reads using the functions *read()* and *getline()*.

```

package inst_buf_array is
    type inst_buf is array (0 to 63) of std_logic_vector(24 downto 0);    -- inst_buff array input
    shared variable sig_inst_buf : inst_buf;                                -- All 64 instructions in instruction buffer

    -- File shenanigans
    file MC_file : text;
    shared variable MC_line : line;
    shared variable MC_string : string(1 to 25);

    -- File procedures
    procedure read_inst;

    -- string to std_logic_vector function
    function string_to_slv (temp_MC_string : string) return std_logic_vector;
end package inst_buf_array;

```

```

package body inst_buf_array is

    function string_to_slv (temp_MC_string : string) return std_logic_vector is
        variable MC_slv : std_logic_vector(24 downto 0);
    begin
        for i in 1 to 25 loop
            if temp_MC_string(i) = '1' then
                MC_slv(24 - (i-1)) := '1';
            else
                MC_slv(24 - (i-1)) := '0';
            end if;
        end loop;
        return MC_slv;
    end function string_to_slv;

    procedure read_inst is
        variable MC_slv_result : std_logic_vector(24 downto 0);
    begin
        file_open(MC_file, "input.txt", read_mode);

        for i in 0 to 63 loop
            readline(MC_file, MC_line);
            read(MC_line, MC_string);
            -- read line type and convert to string

            MC_slv_result := string_to_slv(MC_string);
            -- slv_result from string_to_slv

            sig_inst_buf(i) := MC_slv_result;
        end loop;

        file_close(MC_file);
    end procedure read_inst;

end package body inst_buf_array;

```

4-Stage Pipeline

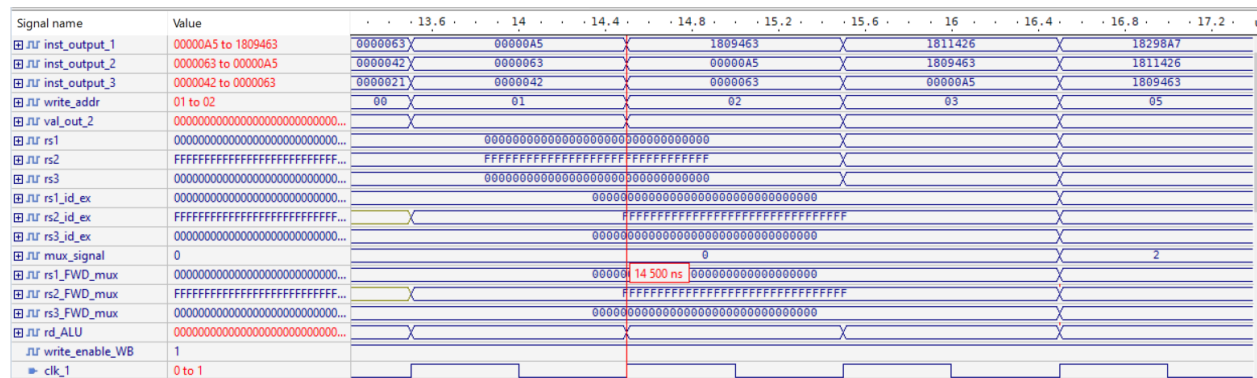
Our processor is meant to be a 4-Stage Pipeline processor, where the four stages are represented by Instruction Fetch (IF), Instruction Decode (ID), Execute (EX), and Write Back (WB). As such, our processor can work on 4 different instructions at the same time, but all at different stages. The dataflow in our processor flows from the IF stage, to the ID stage, then to the EX stage, and finally to the WB stage. (IF -> ID -> EX -> WB) As a design rule, the 4 stages represent the components between the 3 buffer registers (IF/ID Register, ID/EX Register, and the EX/WB Register). These registers are meant to store the inputs to components in the next stage, the outputs of the components from the previous stage, and any control signals or instructions that need to pass on from stage to stage. These registers store their values based on a clock signal, while the components between the registers are not dictated by a clock signal.

In the IF stage, the Instruction Buffer component is there to fetch the current instruction and iterate the Program Counter (PC), and to pass the fetched instruction to the IF/ID Register. In the ID stage, the Register File takes in the instruction from the IF/ID Register and returns the relevant registers from the instruction provided (rs3, rs2, rs1, rd), outputs them to the ID/EX Register, and updates any register specified by the “write_enable” and “write_addr” signals from the WB stage. In the EX stage, both execution of the instruction passed in, and data forwarding happens. First, the instruction passed in from the ID/EX Register and the address of the register to be written to in the Register File (“write_addr” from the WB stage) are passed to the Data Forwarding Unit, where its output is a signal to the Forwarding Muxes (“mux_signal”) that decide

if data forwarding should occur, and to which mux data forwarding should happen. If the “mux_signal” is “00”, then data forwarding does not occur, but if the “mux_signal” was instead “01”, “10”, or “11”, then the values of rs1, rs2, and rs3 respectively should be data-forwarded by the value of the register that is being written to (“val_out”). “Data-forwarded” meaning that the value of “val_out” replaces the value represented in rs1, rs2, or rs3. Even when data forwarding occurs, the Register File would also be updated by the value of “val_out” at the register specified by the register from “write_addr”. Once the outputs to the Forwarding muxes are decided, then they are inputted to the ALU along with the instruction, and compute the correct 128-bit value and output it to the EX/WB Register.

Finally, in the WB stage, depending on the instruction (every instruction except “NOP”) the “write_enable” signal is set, letting the Register File know that a register is going to be updated. The “write_addr” is also produced to tell the Register File and the Data Forwarding Unit which register is being updated along with the value being computed by the ALU (“val_out”) also outputted to both components.

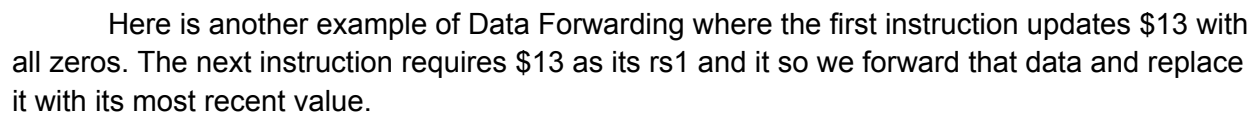
Below is an example from the waveform produced by our simulation showing how 4 instructions are processed in their respective stages:



By looking at the cycle right after the cursor, we see that our processor is able to process 4 instructions all at the same time but at different stages. The signal “inst_output_1” represents the instruction currently fetched by the Instruction Buffer in the IF stage. “inst_output_2” represents the instruction in the ID stage. “inst_output_3” represents the instruction in the EX stage, and “write_addr” represents the register that will be written to in the WB stage. Looking at the cycles immediately after the one pointed by the cursor, we see that between each of the “inst_output” instructions and the “write_addr” signals, the instructions from before are passed through to the next stage and the IF stage fetches a new instruction into the pipeline.


Below are the outputs to the “results.txt” file of the two cycles represented by the waveform that are right after where the cursor is pointing:


```
LI 0, 5, $5
SHRHI $5, $3, $3
```

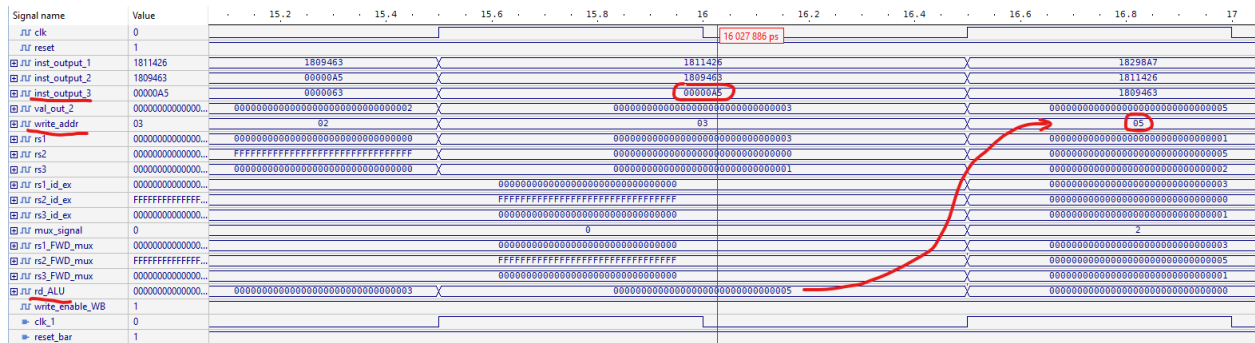


(Instruction rs2, rs1, rd)

```
AND $12, $1, $13
INVB $0, $13, $13
```

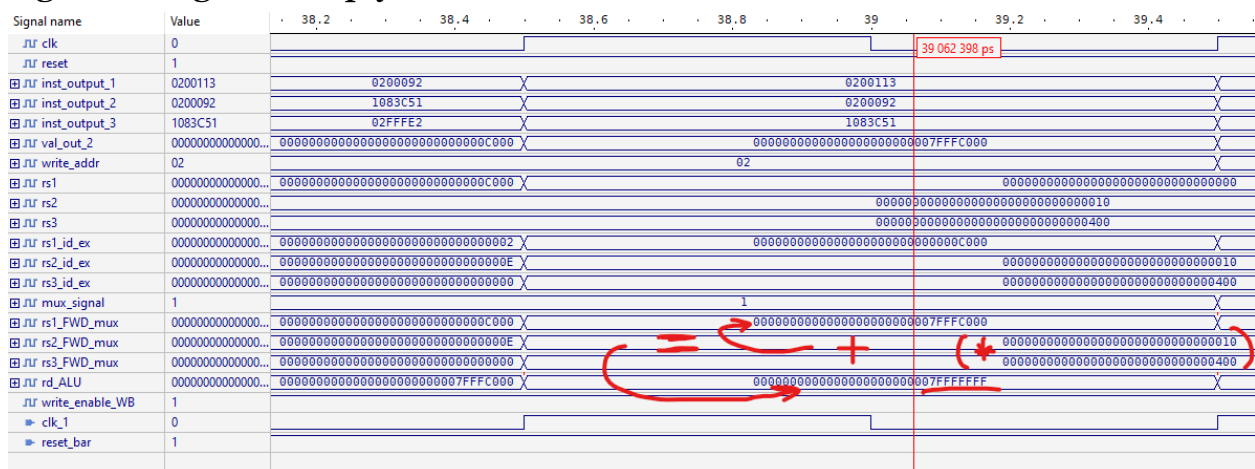


LI



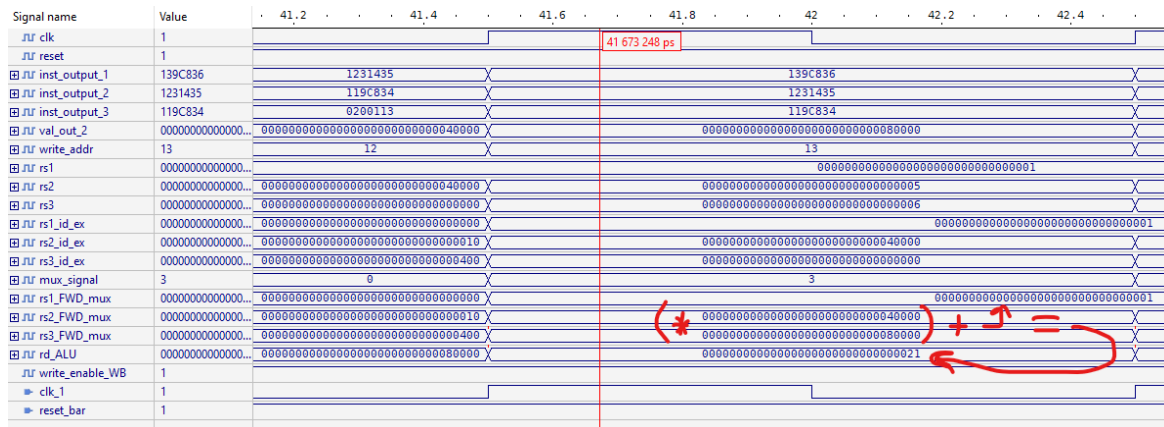
Here we load \$5 as shown by *write_addr* with a value of 5 as shown by *rd_ALU* and the instruction is 00000A5 in hex which is equivalent to LI 0, 5, \$5 (**LI load index, 16-bit immediate, rd**) in our own MIPS syntax.

Signed Integer Multiply-Add Low with Saturation



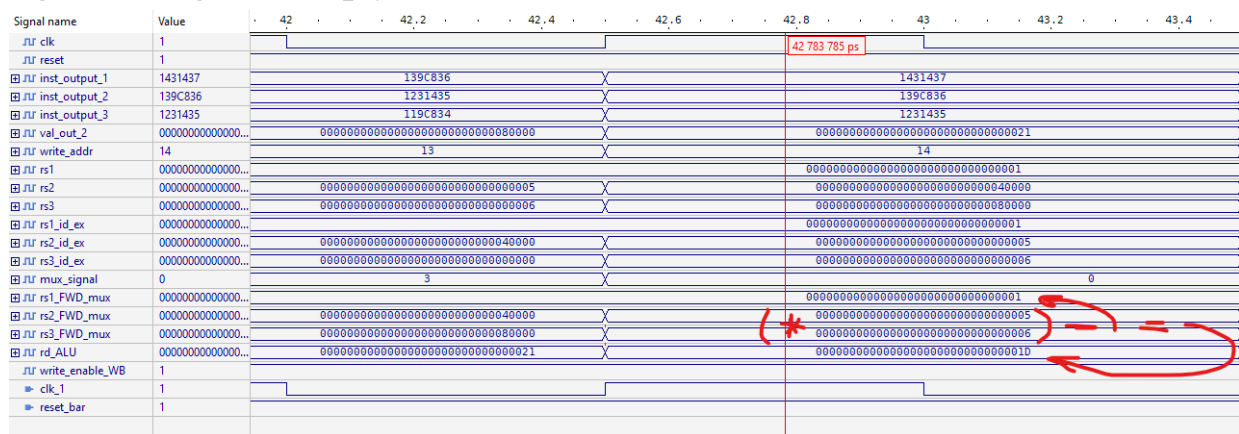
Multiply the low 16 bits of rs3 by the low 16 bits of rs2 and then add rs1 then check for overflow and saturation. In our test we used [31:0] to test our overflow case, by doing $1024(rs3) * 16(rs2) + (2^{17})$ we get overflow and as you can see we get 7FFFFFFF in rd which is the max value we can represent in 32 bits. **(Other cases were shown in the previously submitted report)**

Signed Integer Multiply-Add High with Saturation



Multiply the high 16 bits of rs3 by the high 16 bits of rs2 and then add rs1 then check for overflow and saturation. In our test we used [31:0] to do $8 * 4 + 1 = 33$ which is 21 in hex.

Signed Integer Multiply-Subtract Low with Saturation



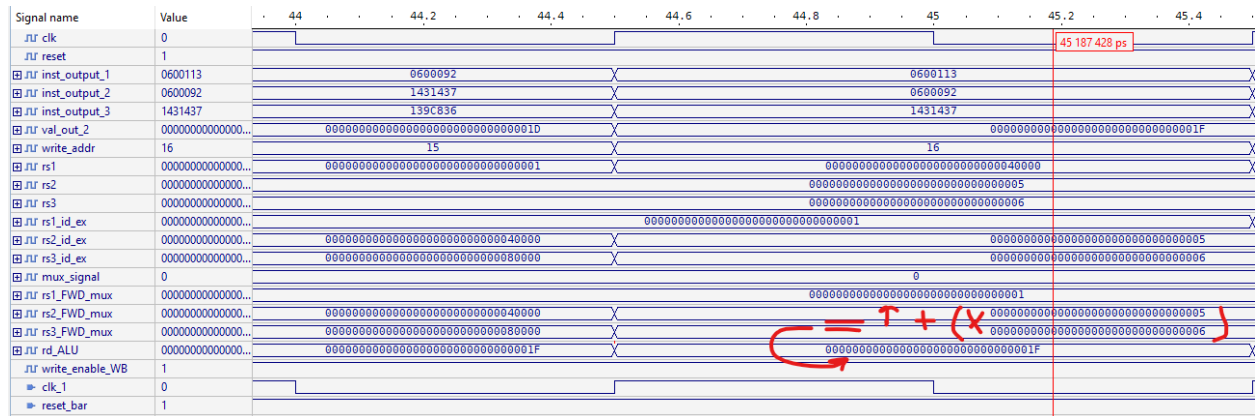
Multiply the low 16 bits of rs3 by the low 16 bits of rs2 and then subtract rs1 then check for overflow and saturation. In our test we used [31:0] to do $6 * 5 - 1$ or 1D in hex.

Signed Integer Multiply-Subtract High with Saturation



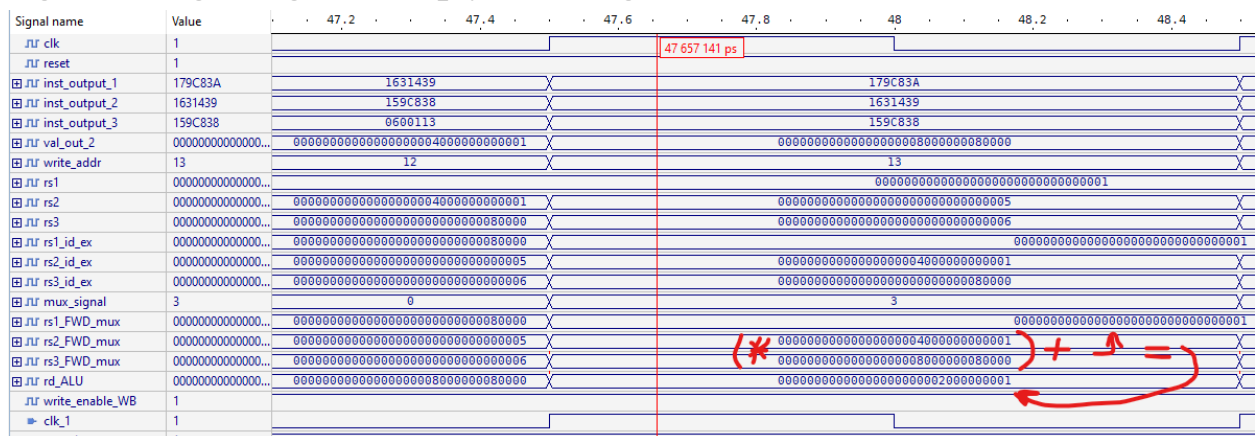
Multiply the high 16 bits of rs3 by the high 16 bits of rs2 and then subtract rs1 then check for overflow and saturation. In our test we used [31:0] to show $8 * 4 - 1 = 31$ or 1F in hex.

Signed Long Integer Multiply-Add Low with Saturation



Multiply the low 32 bits of rs3 by the low 32 bits of rs2 and then add rs1 then check for overflow and saturation. In our test we used [63:0] to show $6 * 5 + 1 = 31$ or 1F in hex.

Signed Long Integer Multiply-Add High with Saturation




Multiply the high 32 bits of rs3 by the high 32 bits of rs2 and then add rs1 then check for overflow and saturation. In our test we used [63:0] to do $524288 * 262144 + 1 = 137438953473$.

Signed Long Integer Multiply-Subtract Low with Saturation




Multiply the low 32 bits of rs3 by the low 32 bits of rs2 and then subtract rs1 then check for overflow and saturation. In our test we used [63:0] to do $6 * 5 - 1 = 29$ or 1D in hex.

Signed Long Integer Multiply-Subtract High with Saturation



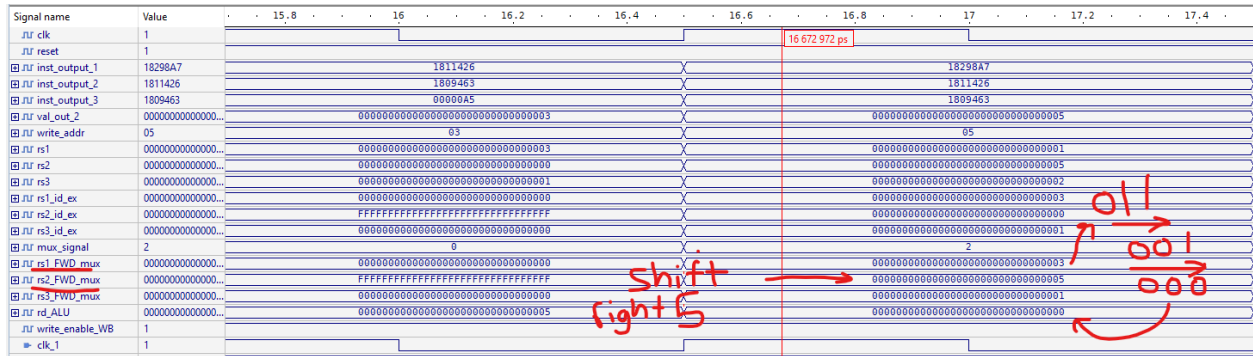
Multiply the high 32 bits of rs3 by the high 32 bits of rs2 and then subtract rs1 then check for overflow and saturation. In our test we used [63:0] to do $524288 * 262144 - 1 = 137438953471$.

NOP



With NOP (inst = 1800000), after taking in values for rs1, rs2, and rs3, the value of rd does not change at all, just as the NOP instruction is supposed to do and our write_enable = 0 so the register file is not tampered with.

SHRHI

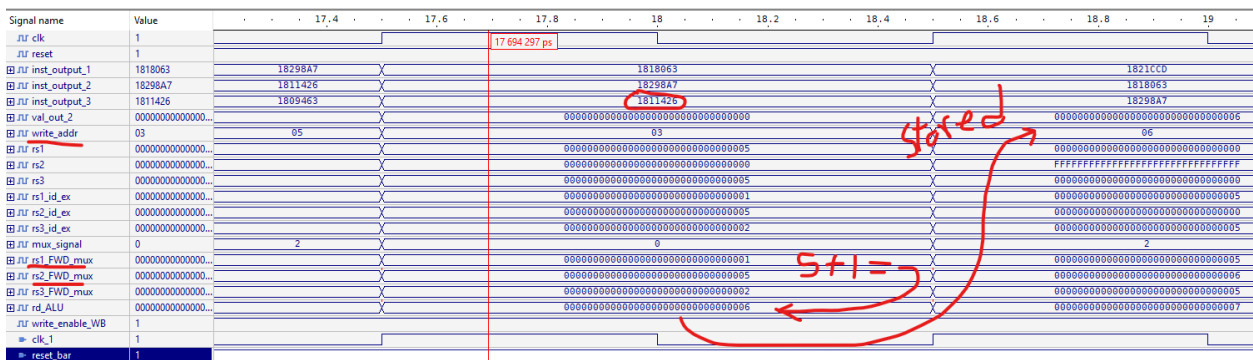


With SHRHI, (inst = 1809463),
(Instruction rs2, rs1, rd)

SHRHI \$5, \$3, \$3

Register \$5 stores a value of 5 and register \$3 stores a value of 3. Here we are shifting \$3 over 5 to the right which then creates a value of 0 which is then stored back into \$3.

AU




With AU, (inst = 1811426),
(Instruction rs2, rs1, rd)

AU \$5, \$1, \$6

The packed 32-bit words in \$5 and \$1 have the values 0, 0, 0, 5 and 0, 0, 0, 1 (from the most significant words to the least significant) respectively. After AU, the correct result of 0, 0, 0, 6 in the packed 32-bit words of \$6 are correctly computed.

CNT1H




With CNT1H, (inst = 1818063),
(Instruction rs2, rs1, rd)

CNT1H \$0, \$3, \$3

The counts of '1' in each halfword in \$3 are computed back into \$3, as the values of 0, 0, 0, 0, 0, 0, 0, and 3 (from the most significant halfwords to the least significant) are represented in the corresponding halfwords of rd which is \$3 as well, with the counts of 0, 0, 0, 0, 0, 0, 0, and 2 respectively.

AHS

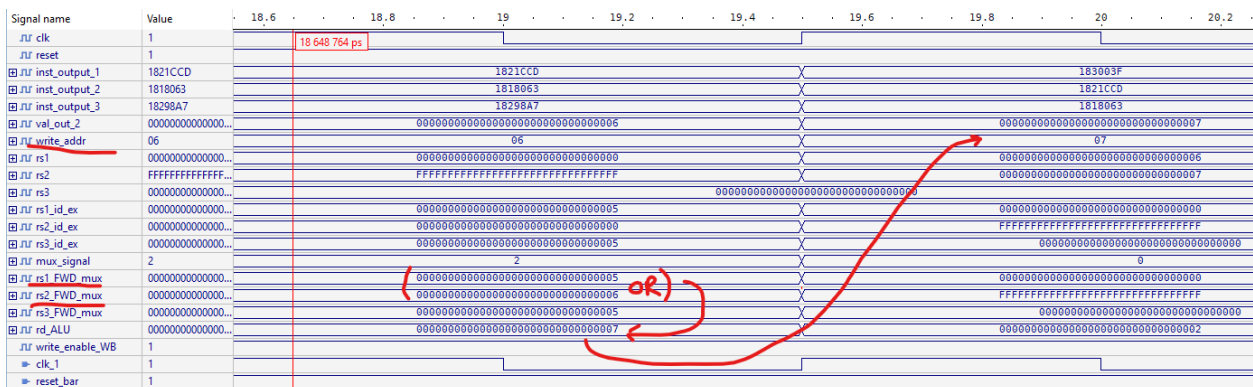


With AHS, (inst = 1821CCD),
(Instruction rs2, rs1, rd)

AHS \$7, \$6, \$13

We test the least significant 16-bits to do 7 + 6 using registers \$7 and \$6 which store those values and get 13 which is then stored into \$13.

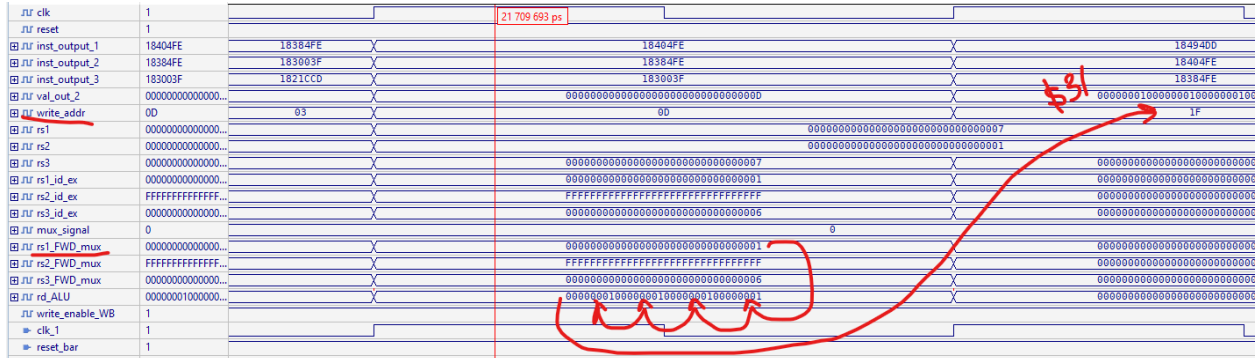
OR



With OR, (inst = 18298A7),
(Instruction rs2, rs1, rd)
 OR \$6, \$5, \$7

The logical OR of \$6 which stores 6 (0110) and \$5 which stores 5 (0101) gets us 7 (0111) and it's stored in \$7.

BCW




With BCW, (inst = 183003F),
(Instruction rs2, rs1, rd)

BCW \$0, \$1, \$31

The least significant word of \$1 (which is 1 represented in 32-bits) is copied to all of the 4 32-bit words of \$31, which is clearly shown correctly in the waveform.

MAXWS



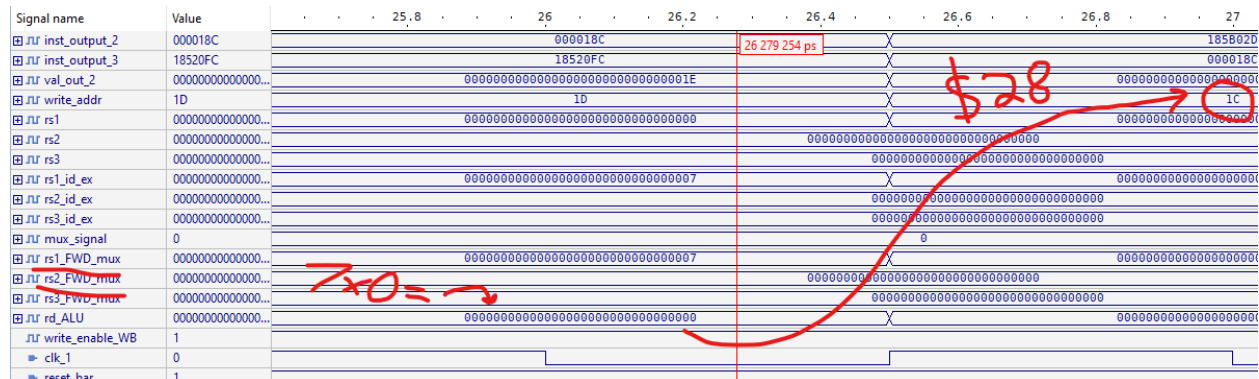
With MAXWS, (inst = 18384FE),
(Instruction rs2, rs1, rd)

MAXWS \$1, \$7, \$30

Out of the values from the 1st word of \$7 (value of 7) and \$1 (value of 1), $7 > 1$ so \$30 correctly picks out the **highest** value in each corresponding word, in this case we only use 1 word, and since 7 is greater than 1, \$30 stores 7 in its first word.

MINWS

MLHSS



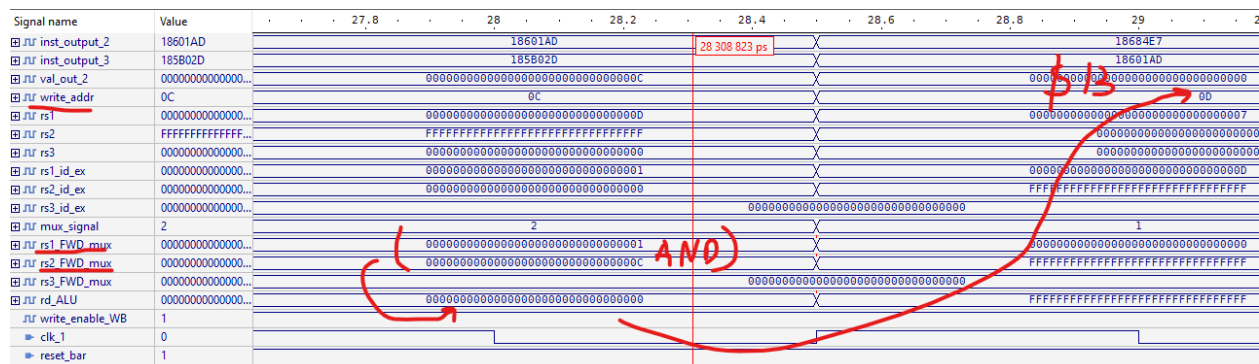
With MLHSS, (inst = 18520FC),

(Instruction rs2, rs1, rd)

MLHSS \$8, \$7, \$28

We are multiplying each halfword of \$7 by the sign of the corresponding halfword in \$8 (in essence taking the 2's complement of the number in \$7, and adjusting it for saturated rounding, and if a halfword in \$8 has the value of 0 (this value returns a result of zero in the corresponding halfword in \$28)). In our case, \$8 is filled with 0 so we have $0 * 7 = 0$ in \$28..

AND



With AND, (inst = 185B02D),

(Instruction rs2, rs1, rd)

AND \$12, \$1, \$13

The logical AND of \$1 with a value of 1 (0001) and \$12 with a value of 12 (1100) is 0 (0000) and this gets stored in \$13

INVB

[illegible]

SFWU \$14, \$10, \$4

The difference of the corresponding words of \$14 from \$10 are outputted to \$4, where in our waveform, we subtract the values 14 - 10 to get an output of 4 in \$4.

[illegible]

SFHS \$14, \$10, \$9
 We do the packed 16-bit halfword subtraction between \$14 and \$10, and use saturation rounding if needed for the difference. Since we only do 14 - 10 saturation is not necessary and the result of 4 is placed into \$9.

Waveform:

Below is a snippet of the waveform from our simulation:


```
1100000111000010011111110
1100001000000010011111110
1100001001001010011011101
1100001010010000011111100
00000000000000000110001100
1100001011011000000101101
```

Binary Machine Code

[illegible]

As mentioned from before, the “IF”, “ID”, “EX”, and “WB” fields represent the instructions present in the IF, ID, and EX stages, along with the address of the register that will be written back to the Register File in the WB stage. “rs1”, “rs2”, and “rs3” fields represent the contents of

the registers specified by rs1, rs2 and rs3 from the instruction in the EX stage. “Rd” represents the output from the ALU. “Mux Signal” represents the mux signal decided by the Data Forwarding Unit to the Forwarding Muxes, “Write Enable” represents the “write_enable” signal from the output of the EX/WB Register, and “Val Out” is the value from the “rd_out” signal from the EX/WB Register.

Going into detail about what we see in the Results File, since all of the instructions in these four cases do not cause any data hazards between themselves, the mux signal is always “00”, signifying that data forwarding was not needed for any of the instructions to be executed. The “write_enable” signal is also always set to ‘1’ since the instructions in the four cycles did not include “NOP” as an instruction. Only the “NOP” instruction is the instruction where the “write_enable” signal is set to ‘0’. Finally, in each of the cycles, the result from the ALU (or the result of the instruction in the EX stage) is represented by the “rd” field, and the result from the EX/WB Register (or the result of the instruction in the WB stage) is represented by the “Val Out” field. **(Since we set the reset signal to be ‘0’ for two cycles, the first cycle of execution from the processor starts at Cycle #3 in the Results File.)**

Expected Results File:

Below is our Expected Results (expected_results.txt) for the values in the 32 registers in the Register File once all of the 64 instructions have run:

Expected Results:

[illegible]

Comparing the results from the Results File to the expected results, we see that from the waveform, only the first 3 MIPS instructions were executed in the EX stage. As such, we can compare the values of the “rd” registers for these 3 instructions with our expected results.

"00
001" from the rd register in Cycle #24. This is the exact same as the value of register \$30 from our expected results file. With the 3rd MIPS instruction ("MLHU \$5, \$6, \$29") that we can see that executed from the waveform, the value of \$29 from the ALU was "00
0011110", which is the same as the our expected value for the value of \$29 at the end of execution.

Conclusion

Working on this project gave us a lot of insight on how pipelining inside a processor actually works. Data forwarding as a concept seems very challenging to understand and gave us a hard time thinking about how to implement it. When we finally realized how to create the unit, it was a lot easier than we initially thought. Despite some challenges, our Processor achieves all the goals that were set according to the results we printed and read. If we had more time we could improve our data efficiency as passing the whole 25-bit instruction through the processor is very inefficient.